

Class 10: Halloween mini project

Fan Wu(PID:A15127541)

Table of contents

Background	1
Data import	1
Quick overview of the dataset	4
Overall Candy Rankings	11
Winpercent and Pricepercent	16
Exploring the correlation structure	19
Principal Component Analysis	20

Background

As it is nearly Halloween and the half way point in the quarter, let's do a mini project to help us figure our the best candy!

Data import

Our data comes from the 538 website and is available as a CSV file:

```
candy <- read.csv("candy-data.csv", row.names = 1)
```

```
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisped	rice	wafer
100 Grand	1	0	1			0	0		1
3 Musketeers	1	0	0			0	1		0
One dime	0	0	0			0	0		0
One quarter	0	0	0			0	0		0
Air Heads	0	1	0			0	0		0

	Almond Joy	1	0	0	1	0	0
		hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173	
3 Musketeers	0	1	0	0.604	0.511	67.60294	
One dime	0	0	0	0.011	0.116	32.26109	
One quarter	0	0	0	0.011	0.511	46.11650	
Air Heads	0	0	0	0.906	0.511	52.34146	
Almond Joy	0	1	0	0.465	0.767	50.34755	

```
flextable::flextable(head(candy, 10))
```

chocolate	fruity	caramel	peanut	almond	nougat	crispedrice	wafer	hard	bar	pluribus	s
1	0	1	0	0	0	1	0	0	1	0	0
1	0	0	0	0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0	0	0	0
1	0	0	1	0	0	0	0	0	1	0	0
1	0	1	1	1	0	0	0	0	1	0	0
0	0	0	1	0	0	0	0	0	0	1	0
0	0	0	0	0	0	0	0	0	0	0	1
0	1	1	0	0	0	0	0	0	0	0	0

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

```
candy |>
  nrow()
```

```
[1] 85
```

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

My favorite winpercent

```
candy["Twix", ]$winpercent
```

```
[1] 81.64291
```

```
## plyr is used to pull out wire by electricians
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
## filter() and select() are useful
```

```
candy |>
  filter(rownames(candy) == "Twix") |>
  select(winpercent)
```

```
      winpercent
Twix    81.64291
```

```
candy |>
  filter(rownames(candy) == "Almond Joy") |>
  select(winpercent)
```

```
    winpercent
Almond Joy    50.34755
```

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

```
candy["Hershey's Special Dark", ]$winpercent
```

```
[1] 59.23612
```

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

Quick overview of the dataset

```
## Only calling on the skim() from the package,
## not loading other stuff from the package
skimr::skim(candy)
```

Table 2: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

Yes, the winpercent column

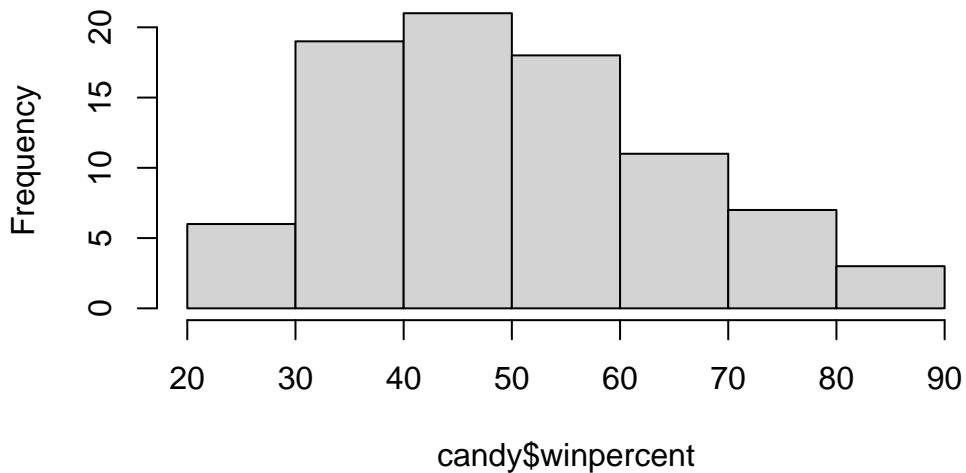
Q7. What do you think a zero and one represent for the candy\$chocolate column?

There's no chocolate in this candy

Q8. Plot a histogram of winpercent values

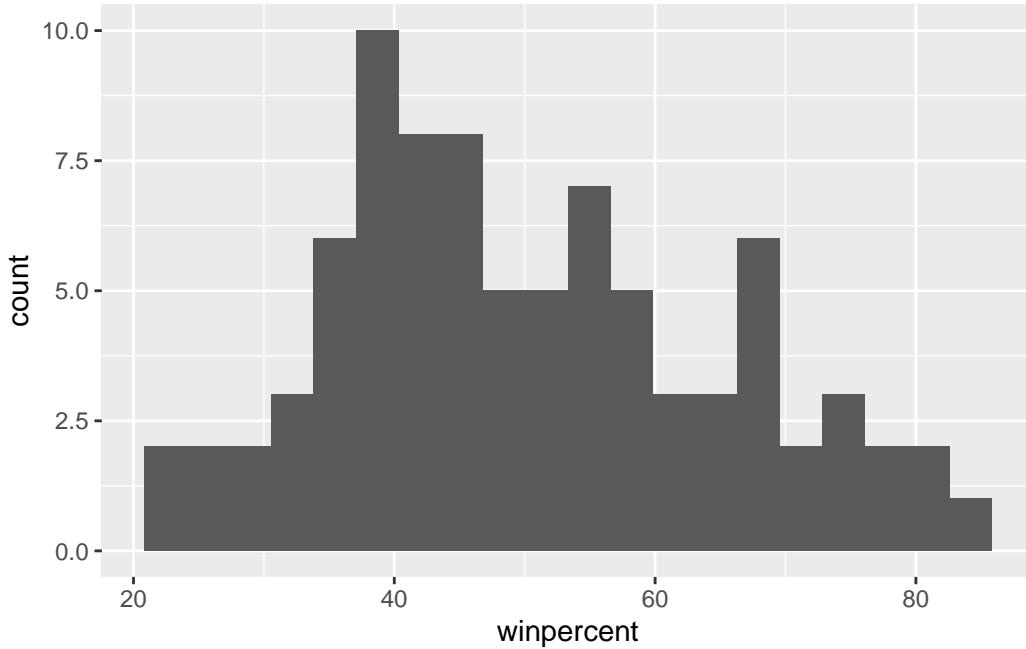
```
hist(candy$winpercent)
```

Histogram of candy\$winpercent



```
library(ggplot2)

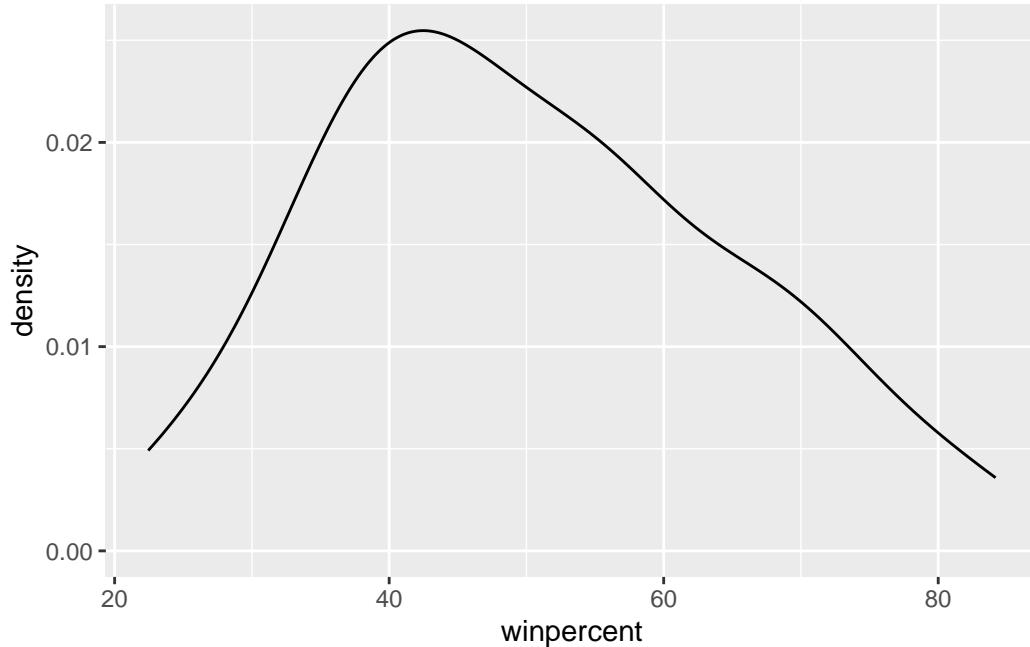
ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins = 20)
```



Q9. Is the distribution of winpercent values symmetrical?

Not symmetrical

```
ggplot(candy) +  
  aes(winpercent) +  
  geom_density()
```



Q10. Is the center of the distribution above or below 50%?

```
mean(candy$winpercent)
```

[1] 50.31676

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

The center of distribution is below 50%, since the median is 47.83

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

On average, chocolate candy is higher ranked than fruit candy

```
#1. Find all chocolate candy in the dataset
#2. Find their winpercent values
#3. Calculate the mean of these values
```

```
#4. Repeat 1,2,3 for fruity candy

#5. mean winpercent of chocolate vs. mean winpercent of fruity

choc inds <- candy$chocolate==1
choc.win <- candy[choc inds, ]$winpercent
choc.mean <- mean(choc.win)

choc.mean
```

[1] 60.92153

```
fruit inds <- candy$fruity==1
fruit.win <- candy[fruit inds, ]$winpercent
fruit.mean <- mean(fruit.win)

fruit.mean
```

[1] 44.11974

```
# try using dplyr

candy |>
  filter(chocolate == 1) |>
  select(winpercent)
```

	winpercent
100 Grand	66.97173
3 Musketeers	67.60294
Almond Joy	50.34755
Baby Ruth	56.91455
Charleston Chew	38.97504
Hershey's Kisses	55.37545
Hershey's Krackel	62.28448
Hershey's Milk Chocolate	56.49050
Hershey's Special Dark	59.23612
Junior Mints	57.21925
Kit Kat	76.76860
Peanut butter M&M's	71.46505
M&M's	66.57458

Milk Duds	55.06407
Milky Way	73.09956
Milky Way Midnight	60.80070
Milky Way Simply Caramel	64.35334
Mounds	47.82975
Mr Good Bar	54.52645
Nestle Butterfinger	70.73564
Nestle Crunch	66.47068
Peanut M&Ms	69.48379
Reese's Miniatures	81.86626
Reese's Peanut Butter cup	84.18029
Reese's pieces	73.43499
Reese's stuffed with pieces	72.88790
Rolo	65.71629
Sixlets	34.72200
Nestle Smarties	37.88719
Snickers	76.67378
Snickers Crisper	59.52925
Tootsie Pop	48.98265
Tootsie Roll Juniors	43.06890
Tootsie Roll Midgies	45.73675
Tootsie Roll Snack Bars	49.65350
Twix	81.64291
Whoppers	49.52411

Q12. Is this difference statistically significant?

Since the p-value is very small, the difference is statistically significant.

```
t.test(choc.win, fruit.win)
```

```
Welch Two Sample t-test

data: choc.win and fruit.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
## using dplyr
candy |>
  arrange(winpercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat		
Nik L Nip	0	1	0	0	0	0		
Boston Baked Beans	0	0	0	1	0	0		
Chiclets	0	1	0	0	0	0		
Super Bubble	0	1	0	0	0	0		
Jawbusters	0	1	0	0	0	0		
	crisped	rice	wafer	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0	1	0	0.197	0.976	
Boston Baked Beans	0	0	0	1	0	0.313	0.511	
Chiclets	0	0	0	1	0	0.046	0.325	
Super Bubble	0	0	0	0	0	0.162	0.116	
Jawbusters	0	1	0	1	0	0.093	0.511	
	winpercent							
Nik L Nip	22.44534							
Boston Baked Beans	23.41782							
Chiclets	24.52499							
Super Bubble	27.30386							
Jawbusters	28.12744							

```
# base R sort() and order()
x<- c(5, 1, 10, 4)
sort(x)
```

```
[1] 1 4 5 10
```

```
order(x)
```

```
[1] 2 4 1 3
```

```
## using base R
ord.ind <- order(candy$winpercent)
head(candy[ord.ind,], n=5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat				
Nik L Nip	0	1	0	0	0	0				
Boston Baked Beans	0	0	0	1	0	0				
Chiclets	0	1	0	0	0	0				
Super Bubble	0	1	0	0	0	0				
Jawbusters	0	1	0	0	0	0				
	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip	0	0	0	1			0.197		0.976	
Boston Baked Beans	0	0	0	1			0.313		0.511	
Chiclets	0	0	0	1			0.046		0.325	
Super Bubble	0	0	0	0			0.162		0.116	
Jawbusters	0	1	0	1			0.093		0.511	
	win	percent								
Nik L Nip	22.44534									
Boston Baked Beans	23.41782									
Chiclets	24.52499									
Super Bubble	27.30386									
Jawbusters	28.12744									

Q14. What are the top 5 all time favorite candy types out of this set?

```
candy |>
  arrange(-winpercent)|>
  head(5)
```

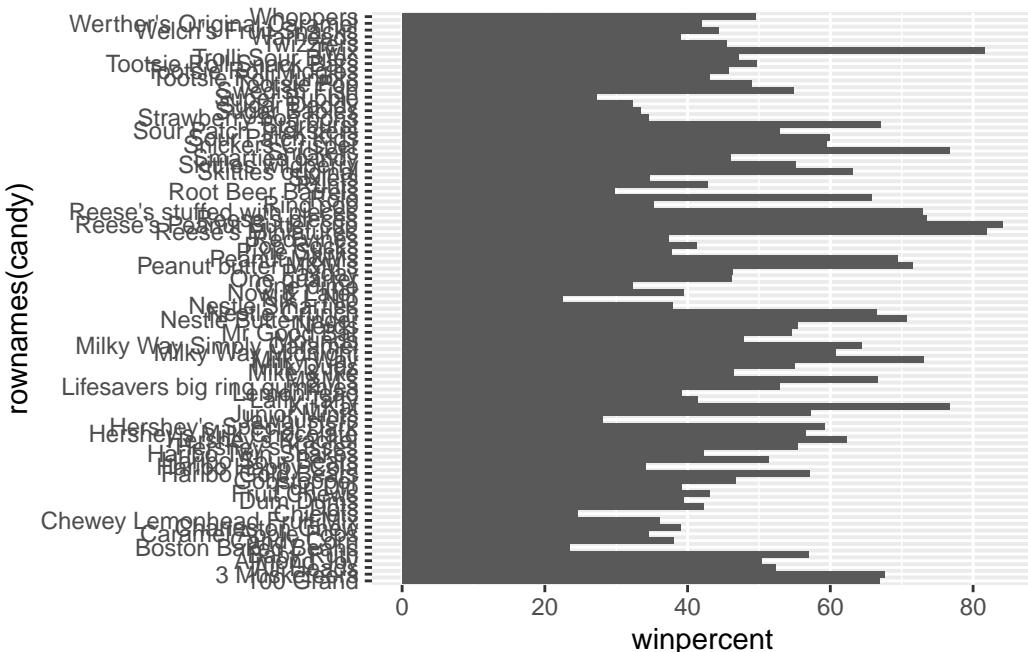
	chocolate	fruity	caramel	peanuty	almondy	nougat			
Reese's Peanut Butter cup	1	0	0	1	0	0			
Reese's Miniatures	1	0	0	1	0	0			
Twix	1	0	1	0	0	0			
Kit Kat	1	0	0	0	0	0			
Snickers	1	0	1	1	1	1			
	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	
Reese's Peanut Butter cup	0	0	0	0			0.720		
Reese's Miniatures	0	0	0	0			0.034		
Twix	1	0	1	0			0.546		
Kit Kat	1	0	1	0			0.313		
Snickers	0	0	1	0			0.546		
	price	percent	win	percent					
Reese's Peanut Butter cup	0.651	84.18029							
Reese's Miniatures	0.279	81.86626							
Twix	0.906	81.64291							
Kit Kat	0.511	76.76860							
Snickers	0.651	76.67378							

```
ord.ind.top <- order(candy$winpercent, decreasing = TRUE)
head(candy[ord.ind.top,], n=5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat		
Reese's Peanut Butter cup	1	0	0		1	0		
Reese's Miniatures	1	0	0		1	0		
Twix	1	0	1		0	0		
Kit Kat	1	0	0		0	0		
Snickers	1	0	1		1	1		
	crisped	rice	wafers	hard	bar	pluribus	sugar	percent
Reese's Peanut Butter cup	0	0	0		0		0.720	
Reese's Miniatures	0	0	0		0		0.034	
Twix	1	0	1		0		0.546	
Kit Kat	1	0	1		0		0.313	
Snickers	0	0	1		0		0.546	
	price	percent	winpercent					
Reese's Peanut Butter cup	0.651	84.18029						
Reese's Miniatures	0.279	81.86626						
Twix	0.906	81.64291						
Kit Kat	0.511	76.76860						
Snickers	0.651	76.67378						

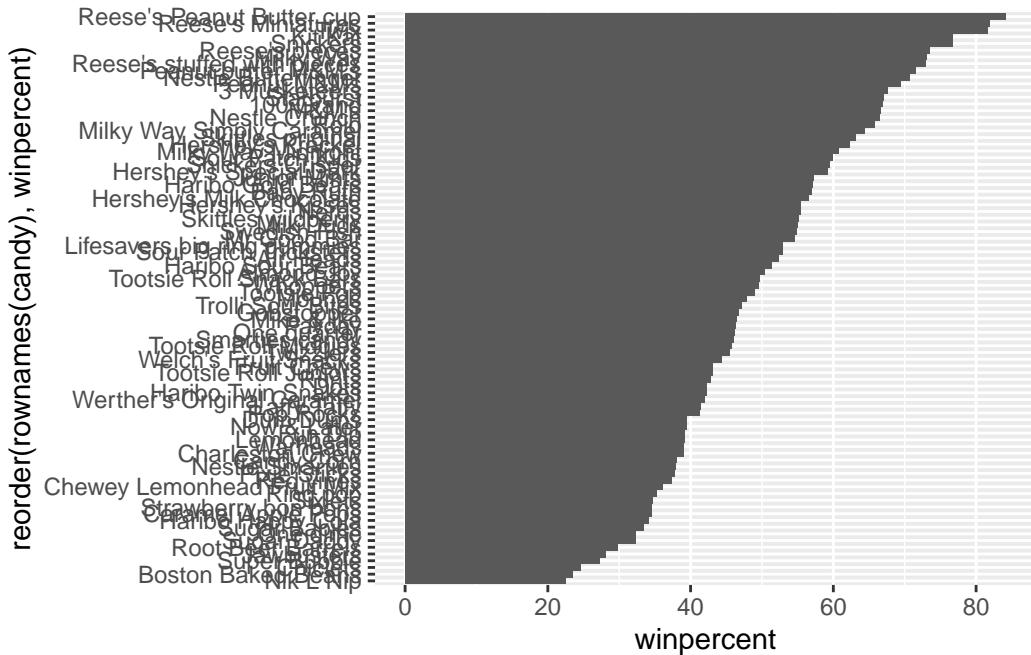
Q15 Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy) +
  aes(x = winpercent,
      y = reorder(rownames(candy), winpercent)) +
  geom_col()
```

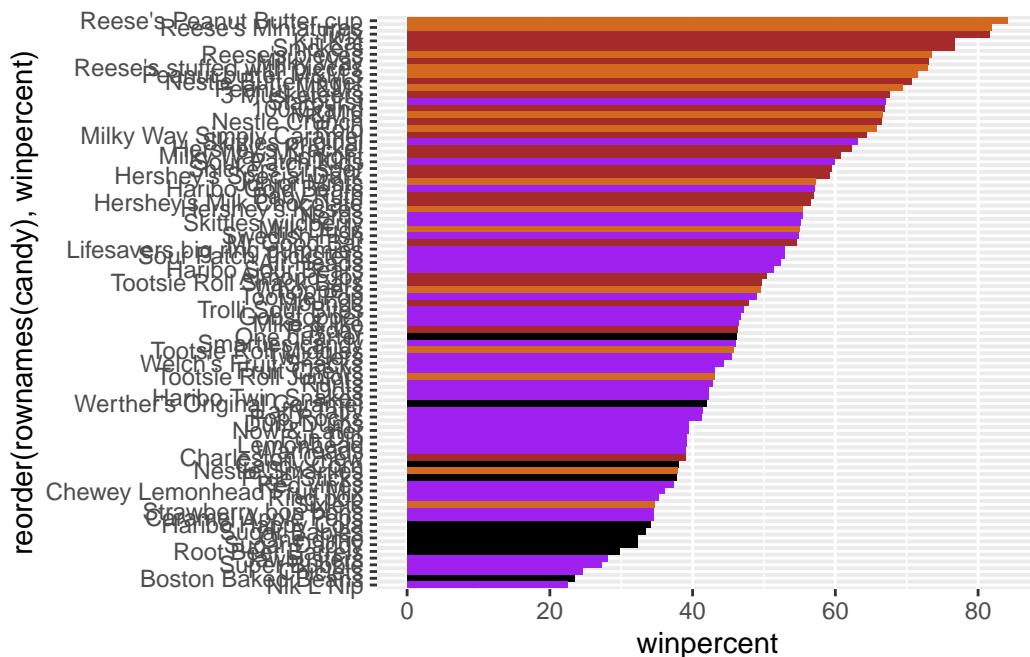


Add some color based on the “type of candy”

```
my_cols <- rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] <- "chocolate"
my_cols[as.logical(candy$fruity)] <- "purple"
my_cols[as.logical(candy$bar)] <- "brown"
my_cols
```

[1]	"brown"	"brown"	"black"	"black"	"purple"	"brown"
[7]	"brown"	"black"	"black"	"purple"	"brown"	"purple"
[13]	"purple"	"purple"	"purple"	"purple"	"purple"	"purple"
[19]	"purple"	"black"	"purple"	"purple"	"chocolate"	"brown"
[25]	"brown"	"brown"	"purple"	"chocolate"	"brown"	"purple"
[31]	"purple"	"purple"	"chocolate"	"chocolate"	"purple"	"chocolate"
[37]	"brown"	"brown"	"brown"	"brown"	"brown"	"purple"
[43]	"brown"	"brown"	"purple"	"purple"	"brown"	"chocolate"
[49]	"black"	"purple"	"purple"	"chocolate"	"chocolate"	"chocolate"
[55]	"chocolate"	"purple"	"chocolate"	"black"	"purple"	"chocolate"
[61]	"purple"	"purple"	"chocolate"	"purple"	"brown"	"brown"
[67]	"purple"	"purple"	"purple"	"purple"	"black"	"black"
[73]	"purple"	"purple"	"purple"	"chocolate"	"chocolate"	"brown"
[79]	"purple"	"brown"	"purple"	"purple"	"purple"	"black"
[85]	"chocolate"					

```
ggplot(candy) +
  aes(x = winpercent,
      y = reorder(rownames(candy), winpercent)) +
  geom_col(fill = my_cols)
```

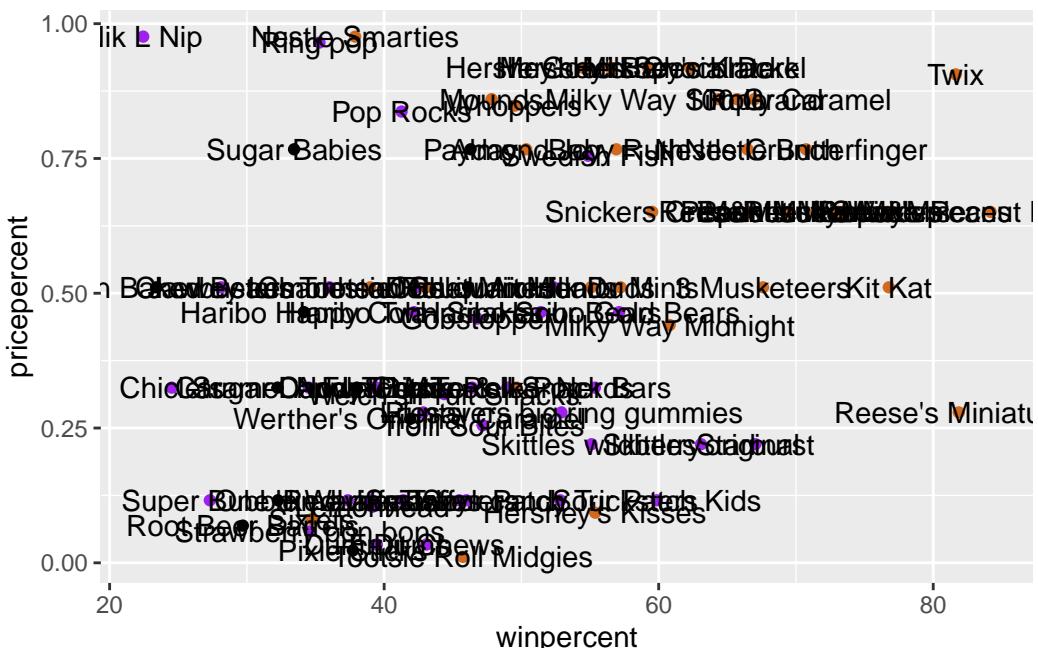


Winpercent and Pricepercent

A plot with both variables/columns winpercent and pricepercent

```
my_points <- rep("black", nrow(candy))
my_points[as.logical(candy$chocolate)] <- "chocolate"
my_points[as.logical(candy$fruity)] <- "purple"
```

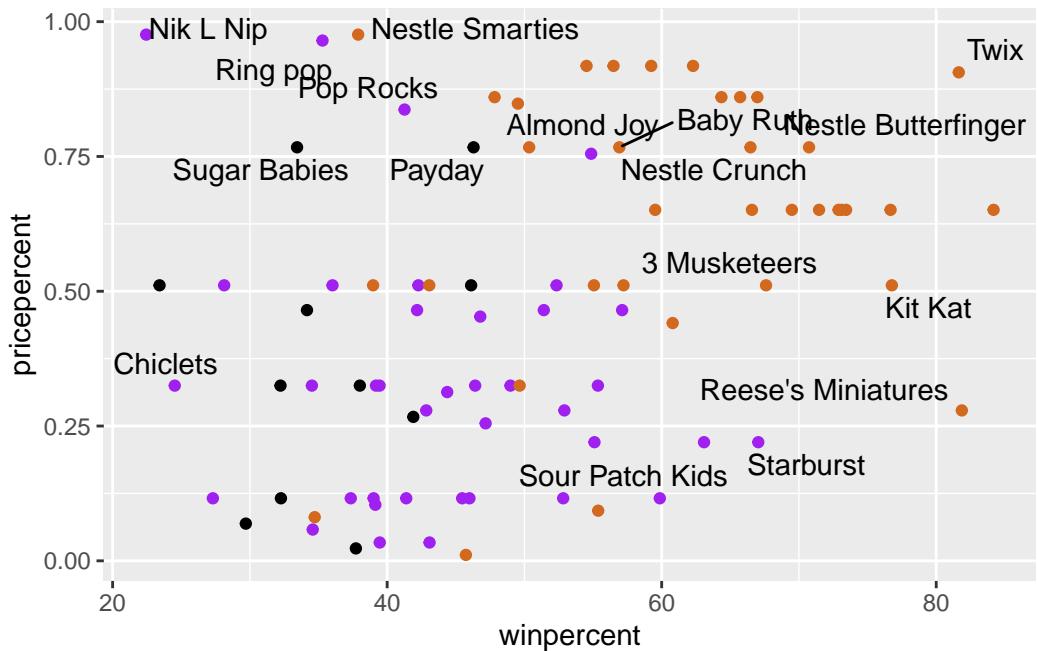
```
ggplot(candy) +
  aes(x = winpercent,
      y = pricepercent,
      label = rownames(candy)) +
  geom_point(col = my_points) +
  geom_text()
```



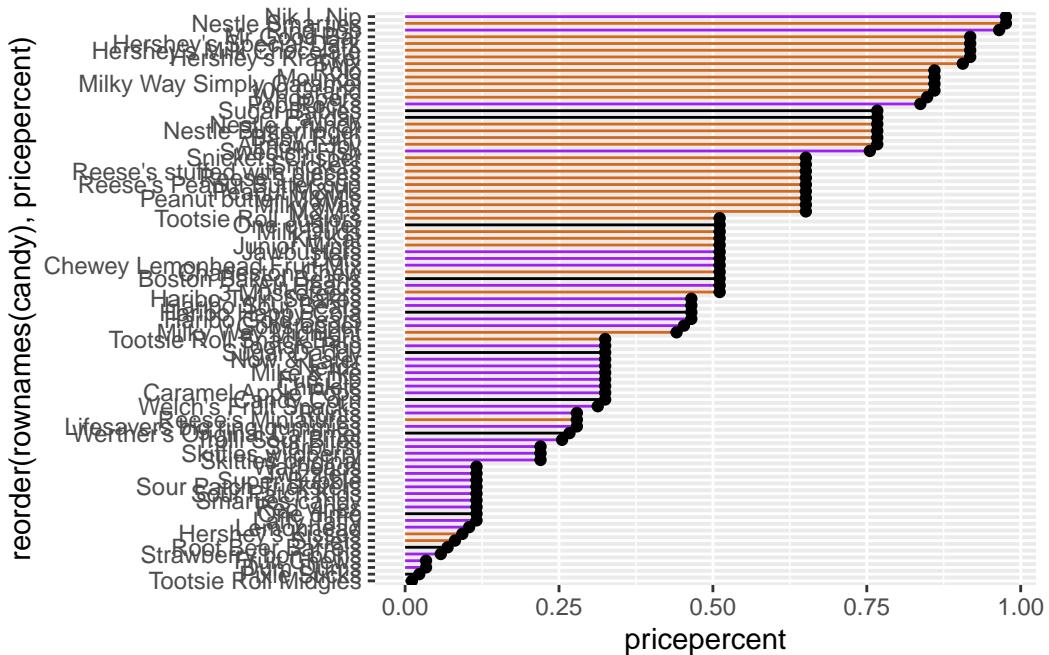
```
library(ggrepel)

ggplot(candy) +
  aes(x = winpercent,
      y = pricepercent,
      label = rownames(candy)) +
  geom_point(col = my_points) +
  geom_text_repel(max.overlaps = 7)
```

Warning: ggrepel: 68 unlabeled data points (too many overlaps). Consider increasing max.overlaps



```
# Make a lollipop chart of pricepercent
ggplot(candy) +
  aes(pricepercent, reorder(rownames(candy), pricepercent)) +
  geom_segment(aes(yend = reorder(rownames(candy), pricepercent),
                    xend = 0), col=my_points) +
  geom_point()
```



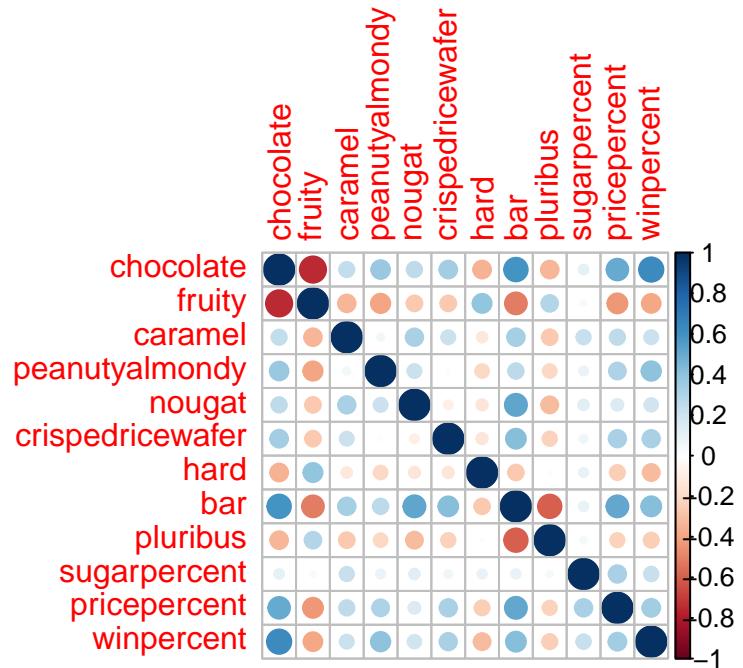
Exploring the correlation structure

Now that we've explored the dataset a little, we'll see how the variables interact with one another. We'll use correlation and view the results with the corrplot package to plot a correlation matrix.

```
# correlation
library(corrplot)

corrplot 0.95 loaded

cij <- cor(candy)
# 1 in cor() result means correlated, negative number results are negatively correlated
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

chocolate and fruity

Q23. Similarly, what two variables are most positively correlated?

chocolate and winpercent

Principal Component Analysis

The function to use is called `prcomp()` w/ an optional `scale=T/F` argument

```
pca <- prcomp(candy, scale = TRUE)
summary(pca)
```

Importance of components:

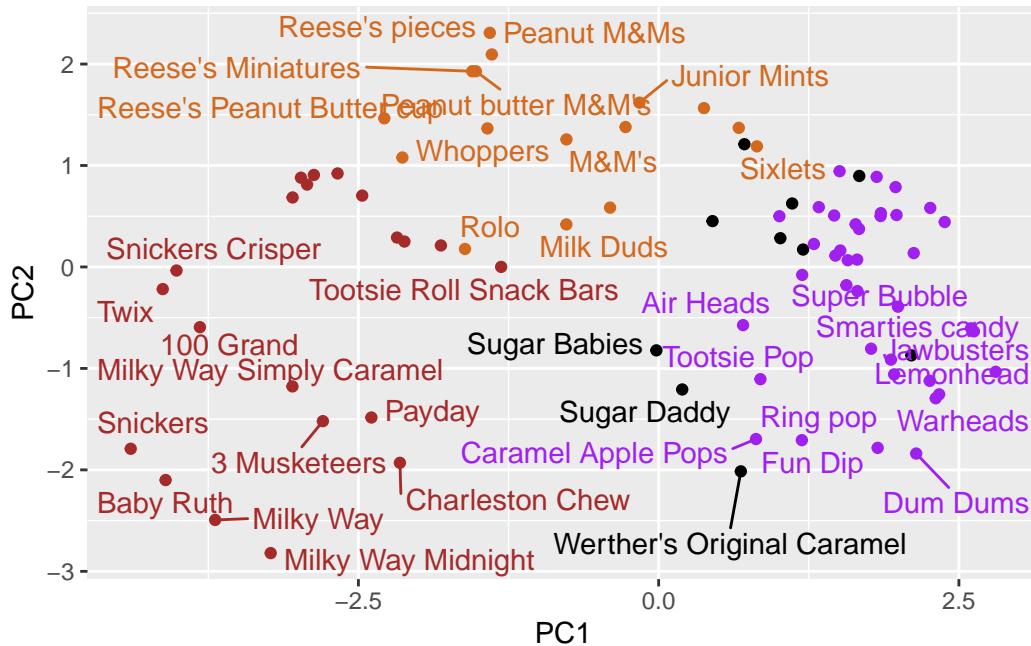
	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		

```
Proportion of Variance 0.04629 0.03833 0.03239 0.01611 0.01317
Cumulative Proportion 0.89998 0.93832 0.97071 0.98683 1.00000
```

Our main PCA result figure

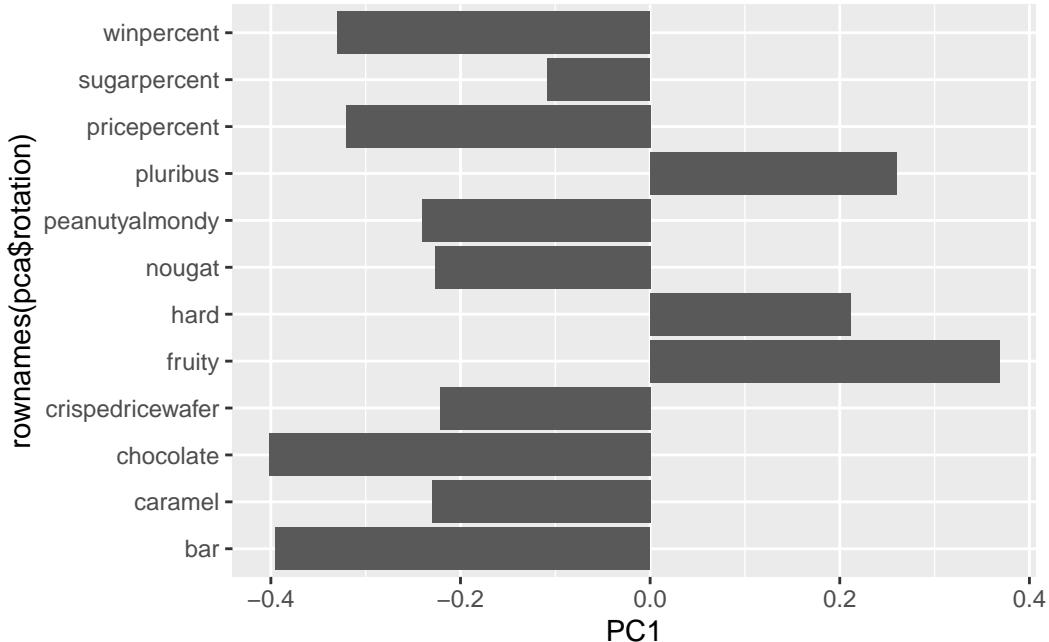
```
ggplot(pca$x) +
  aes(PC1, PC2, label = rownames(pca$x)) +
  geom_point(col = my_cols) +
  geom_text_repel(col = my_cols)
```

Warning: ggrepel: 48 unlabeled data points (too many overlaps). Consider increasing max.overlaps



We should also examine the variable “loadings” or contributions of the original variables to the new PCs

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation))+
  geom_col()
```



```
p <- ggplot(pca$x) +
  aes(PC1, PC2, label = rownames(pca$x)) +
  geom_point(col = my_cols) +
  geom_text_repel(col = my_cols)
```

Interactive plots that can be zoomed on and “brushed” over can be made with the **plotly** package. It’s output is interactive and will not render to PDF :-(

```
library(plotly)
```

```
Attaching package: 'plotly'
```

```
The following object is masked from 'package:ggplot2':
```

```
last_plot
```

```
The following object is masked from 'package:stats':
```

```
filter
```

The following object is masked from 'package:graphics':

layout

```
# ggplotly(p)
```