

Class 12: RNASeq analysis

Fan Wu(PID: A15127541)

Table of contents

Background	1
Data Import	1
Toy differential gene expression	3
DESeq2 analysis	8
Volcano Plot	9
Save our results	10
Add gene annotation	10
Pathway analysis	13
Save our main results	15

Background

Today we will analyze some RNASeq data from Himes et al. on the effects of a common steroid(dexamethasone) on airway smooth muscle cells (ASM cells).

Our starting point is the “counts” data and “metadata” that contain the count values for each gene in their different experiments (i.e. cell lines with or without the drug).

Data Import

```
# Complete the missing code
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

Let's have a wee peak at these objects:

```
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG000000000003	723	486	904	445	1170
ENSG000000000005	0	0	0	0	0
ENSG00000000419	467	523	616	371	582
ENSG00000000457	347	258	364	237	318
ENSG00000000460	96	81	73	66	118
ENSG00000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		
ENSG000000000003	1097	806	604		
ENSG000000000005	0	0	0		
ENSG00000000419	781	417	509		
ENSG00000000457	447	330	324		
ENSG00000000460	94	102	74		
ENSG00000000938	0	0	0		

Q1. How many genes are in this dataset?

```
nrow(counts)
```

```
[1] 38694
```

Q. How many different experiments (column in counts or rows in metadata) are there?

```
ncol(counts)
```

```
[1] 8
```

```
nrow(metadata)
```

```
[1] 8
```

Q2. How many ‘control’ cell lines do we have?

```
table(metadata$dex)
```

control	treated
4	4

```
sum(metadata$dex == "control")
```

```
[1] 4
```

Toy differential gene expression

To start our analysis, let's calculate the mean counts for all genes in the “control” experiments.

1. Extract all “control” columns from the `counts` object
2. Calculate the mean for all rows (i.e. genes) of these “control” column 3-4 Do the same for “treated”
3. compare these `control.mean` and `treated.mean` values

```
#Step 1
control inds <- metadata$dex == "control"
# control.inds <- grep("control", metadata$dex)
control.counts <- counts[, control.inds]

dim(control.counts)
```

```
[1] 38694      4
```

Q3. How would you make the above code in either approach more robust? Is there a function that could help here?

```
rowMeans()
```

```
# calculate mean of each row, i.e. gene, in "control"
control.means <- rowMeans(control.counts)
```

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called `treated.mean`)

```
# Step 3-4
treated.inds <- metadata$dex == "treated"

treated.counts <- counts[, treated.inds]

treated.means <- rowMeans(treated.counts)
```

Store these together for each of bookkeeping as `meanscounts`

```
meanscounts <- data.frame(control.means, treated.means)
head(meanscounts)
```

	control.means	treated.means
ENSG000000000003	900.75	658.00
ENSG000000000005	0.00	0.00
ENSG00000000419	520.50	546.00
ENSG00000000457	339.75	316.50
ENSG00000000460	97.25	78.75
ENSG00000000938	0.75	0.00


Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples. Your plot should look something like the following.

Make a plot of control vs treated mean values for all genes plot per point per gene

```
library(ggplot2)
ggplot(meanscounts, aes(x= control.means, y = treated.means)) +
  geom_point()+
  scale_x_log10() +
  scale_y_log10()
```

Warning in `scale_x_log10()`: log-10 transformation introduced infinite values.

Warning in `scale_y_log10()`: log-10 transformation introduced infinite values.




Make this a log plot

```
plot(meanscounts, log = "xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted from logarithmic plot



We often talk about metrics like “log2 fold-change”

```
# control / treated
log2(10 / 10)
```

```
[1] 0
```

```
# log2(1) = 0, meaning no change, so `log2()` is useful in showing change
```

```
log2(10/20)
```

```
[1] -1
```

```
log2(20/10)
```

```
[1] 1
```

```
log2(40/10)
```

```
[1] 2
```

```
log2(10/40)
```

```
[1] -2
```

Let's calculate the **log2 fold-change** for our treated over control

```
meanscounts$log2fc <-  
  log2(meanscounts$treated.means /  
    meanscounts$control.means)
```

```
head(meanscounts)
```

	control.means	treated.means	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG00000000419	520.50	546.00	0.06900279
ENSG00000000457	339.75	316.50	-0.10226805
ENSG00000000460	97.25	78.75	-0.30441833
ENSG00000000938	0.75	0.00	-Inf

A common “rule of thumb” is a log2 fold-change cutoff of +2 and -2 call genes “Up regulated” or “Down regulated”, when $\log_2(x/y) = +2$ means 4 times difference

Number of upregulated genes

```
sum(meanscounts$log2fc >= 2, na.rm= T)
```

```
[1] 1910
```

Number of downregulated genes at -2 threshold

```
sum(meanscounts$log2fc <= -2, na.rm= T)
```

```
[1] 2330
```

The above data is missing the **statistical significance** evaluation of these statistics: are these changes by chance or are they caused by the drug

DESeq2 analysis

Let's do this analysis properly and keep our inner stats nerd happy - i.e. are the differences we see btw/ drug and no drug significant, given the replicate experiments.

```
library(DESeq2)
```

For DESeq analysis, we need 3 things

1. count values (`countData`)
2. metadata telling us about the columns in `countData` (`colData`)
3. design of the experiment (i.e. what do you want to compare)

Our first function from DESeq2 will setup the input required for analysis by storing all these 3 things together.

```
dds <- DESeqDataSetFromMatrix(countData = counts,
                               colData = metadata,
                               design = ~dex)
```

converting counts to integer mode

```
Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
design formula are characters, converting to factors
```

```
# design = ~dex means this is what we are comparing
```

The main function in DESeq2 that runs the analysis is called `DESeq()`

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

```
final dispersion estimates
```

```
fitting model and testing
```

```
res <- results(dds)
head(res)
```

```
log2 fold change (MLE): dex treated vs control
```

```
Wald test p-value: dex treated vs control
```

```
DataFrame with 6 rows and 6 columns
```

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.194195	-0.3507030	0.168246	-2.084470	0.0371175
ENSG000000000005	0.000000	NA	NA	NA	NA
ENSG00000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026
ENSG00000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106
ENSG00000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691
ENSG00000000938	0.319167	-1.7322890	3.493601	-0.495846	0.6200029

	padj
	<numeric>
ENSG000000000003	0.163035
ENSG000000000005	NA
ENSG00000000419	0.176032
ENSG00000000457	0.961694
ENSG00000000460	0.815849
ENSG00000000938	NA


```
36000 * 0.05
```

```
[1] 1800
```

Volcano Plot

This is a common summary result figure from these types of experiments and plot the **log2 fold-change vs the adjusted p-value**.

```
plot(res$log2FoldChange, -log(res$padj))
abline(v = c(-2,2), col = "red")
abline(h = -log(0.05), col = "red")
```



```
log(0.1)
```

```
[1] -2.302585
```

```
log(0.001)
```

```
[1] -6.907755
```

Save our results

```
write.csv(res, file = "my_results.csv")
```

Add gene annotation

To help make sense of our results, and communicate them to other folks we need to add some more annotation to our main `res` object

We will use two bioconductor packages to first map IDs to different formats, including the classic gene “symbol” gene name.

I will install these with the following commands: `BiocManager::install("AnnotationDbi")`
`BiocManager::install("org.Hs.eg.db")`

```
library(AnnotationDbi)
library(org.Hs.eg.db)
```

Let's see what is in `org.Hs.eg.db` w/ the `columns()` function:

```
columns(org.Hs.eg.db)
```

```
[1] "ACCCNUM"      "ALIAS"        "ENSEMBL"       "ENSEMBLPROT"   "ENSEMBLTRANS"
[6] "ENTREZID"     "ENZYME"       "EVIDENCE"      "EVIDENCEALL"   "GENENAME"
[11] "GENETYPE"     "GO"           "GOALL"         "IPI"          "MAP"
[16] "OMIM"          "ONTOLOGY"     "ONTOLOGYALL"  "PATH"         "PFAM"
[21] "PMID"          "PROSITE"      "REFSEQ"        "SYMBOL"       "UCSCKG"
[26] "UNIPROT"
```

We can translate or “map” IDs between any of these 26 databases using the `mapIds()` function

```
res$symbol <- mapIds(keys = row.names(res),# current IDs
                      keytype = "ENSEMBL", # the format of our IDs
                      x = org.Hs.eg.db ,# where to get mappings from
                      column = "SYMBOL" # the format/DB to map to
)
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 7 columns
  baseMean log2FoldChange      lfcSE      stat    pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.3507030  0.168246 -2.084470 0.0371175
ENSG000000000005  0.000000      NA        NA        NA        NA
```

```

ENSG00000000419 520.134160      0.2061078  0.101059  2.039475 0.0414026
ENSG00000000457 322.664844      0.0245269  0.145145  0.168982 0.8658106
ENSG00000000460 87.682625      -0.1471420  0.257007 -0.572521 0.5669691
ENSG00000000938 0.319167      -1.7322890  3.493601 -0.495846 0.6200029
          padj      symbol
          <numeric> <character>
ENSG00000000003 0.163035      TSPAN6
ENSG00000000005   NA        TNMD
ENSG00000000419 0.176032      DPM1
ENSG00000000457 0.961694      SCYL3
ENSG00000000460 0.815849      FIRRM
ENSG00000000938   NA        FGR

```

Add the mappings for “GENENAME” and “ENTREZID” (NCBI identifier), a and store as `res$genename` and `res$entrez`

```

res$genename <- mapIds(keys = row.names(res),# current IDs
                      keytype = "ENSEMBL", # the format of our IDs
                      x = org.Hs.eg.db ,# where to get mappings from
                      column = "GENENAME" # the format/DB to map to
)

```

`'select()'` returned 1:many mapping between keys and columns

```

res$entrez <- mapIds(keys = row.names(res),# current IDs
                      keytype = "ENSEMBL", # the format of our IDs
                      x = org.Hs.eg.db ,# where to get mappings from
                      column = "ENTREZID" # the format/DB to map to
)

```

`'select()'` returned 1:many mapping between keys and columns

```
head(res)
```

```

log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 9 columns
      baseMean log2FoldChange      lfcSE      stat     pvalue
      <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG00000000003 747.194195      -0.3507030  0.168246 -2.084470 0.0371175

```

ENSG000000000005	0.000000	NA	NA	NA	NA
ENSG000000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026
ENSG000000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106
ENSG000000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691
ENSG000000000938	0.319167	-1.7322890	3.493601	-0.495846	0.6200029
	padj	symbol		genename	entrez
	<numeric>	<character>		<character>	<character>
ENSG000000000003	0.163035	TSPAN6	tetraspanin 6		7105
ENSG000000000005	NA	TNMD	tenomodulin		64102
ENSG000000000419	0.176032	DPM1	dolichyl-phosphate m..		8813
ENSG000000000457	0.961694	SCYL3	SCY1 like pseudokina..		57147
ENSG000000000460	0.815849	FIRRM	FIGNL1 interacting r..		55732
ENSG000000000938	NA	FGR	FGR proto-oncogene, ..		2268

Pathway analysis

Depend on the biological question we are interested in, we will use different functional set databases

Take our different expressed genes, check for overlap of these genes w/ different pathway

There are lots of bioconductor packages to do this type of analysis, now let's just try one called **GAGE** again we need to install this if we don't have it already

```
library(gage)
library(gageData)
library(pathview)
```

To use **GAGE** I need two things

- a named vector of fold-change values for our DEGs (our geneset of interest)
- a set of pathways or genesets to use for annotation.

```
x = c("david" = 5, "barry" = 10)
x
```

```
david barry
5     10
```

```
names(x)
```

```
[1] "david" "barry"
```

```

names(x) <- c("low", "high")
x

  low high
  5   10

foldchanges <- res$log2FoldChange
#name by entrezid, as KEGG only speak in entrezid, not genename
names(foldchanges) <- res$entrez
head(foldchanges)

```

	7105	64102	8813	57147	55732	2268
	-0.35070302	NA	0.20610777	0.02452695	-0.14714205	-1.73228897

```

#run gageData for pathway
data("kegg.sets.hs")

# Get the results
keggres = gage(foldchanges, gsets=kegg.sets.hs)

```

In our results object we have:

```
attributes(keggres)
```

```
$names
[1] "greater" "less"     "stats"
```

```
head(keggres$less, 5)
```

	p.geomean	stat.mean
hsa05332 Graft-versus-host disease	0.0004250461	-3.473346
hsa04940 Type I diabetes mellitus	0.0017820293	-3.002352
hsa05310 Asthma	0.0020045888	-3.009050
hsa04672 Intestinal immune network for IgA production	0.0060434515	-2.560547
hsa05330 Allograft rejection	0.0073678825	-2.501419
	p.val	q.val
hsa05332 Graft-versus-host disease	0.0004250461	0.09053483
hsa04940 Type I diabetes mellitus	0.0017820293	0.14232581
hsa05310 Asthma	0.0020045888	0.14232581

hsa04672 Intestinal immune network for IgA production	0.0060434515	0.31387180
hsa05330 Allograft rejection	0.0073678825	0.31387180
	set.size	exp1
hsa05332 Graft-versus-host disease	40	0.0004250461
hsa04940 Type I diabetes mellitus	42	0.0017820293
hsa05310 Asthma	29	0.0020045888
hsa04672 Intestinal immune network for IgA production	47	0.0060434515
hsa05330 Allograft rejection	36	0.0073678825

Let's look at one of these pathways w/ our genes colored up so we can see the overlap


```
pathview(pathway.id = "hsa05310" ,gene.data = foldchanges)
```

```
'select()' returned 1:1 mapping between keys and columns
```

```
Info: Working in directory /Users/fanwu/Desktop/BIMM143/FA25_Class12
```

```
Info: Writing image file hsa05310.pathview.png
```

Add this pathway figure to our lab report



Save our main results

```
write.csv(res, file = "myresults_annotation.csv")
```